

STAGE 1 *base classification*

Input
(q, r)
question, response

Pass 1 · Base classifier f_θ
($\hat{y}_1, c$) = $f_\theta(q, r)$
 $\hat{y}_1 \in \mathcal{Y} = \{\text{CR}, \text{AMB}, \text{CNR}\}$
verbalized confidence $c \in [0, 1]$

taxonomy
in prompt

structured
output

Evasion taxonomy (clarity map $\kappa : \mathcal{E} \rightarrow \mathcal{Y}$)

CR = { Explicit }

AMB = { Implicit, General, Partial, Dodging, Deflection }

CNR = { Declining, Claims Ignorance, Clarification }

Pass-1 output (JSON)

```
{
  "classification":
  "CR|AMB|CNR",
  "confidence": 0.0-1.0,
  "reasoning": "...
}
```

STAGE 2 *gate*

$c \geq \tau$?

$\tau = 0.95$

$c \geq \tau$: exit early ($\hat{y} = \hat{y}_1$)

162 pairs (52.6% of D_{dev})

AMB ($\hat{y}_1 = \text{AMB}$)

Output
 $\hat{y} \in \mathcal{Y} = \{\text{CR}, \text{AMB}, \text{CNR}\}$
✓

STAGE 3 *asymmetric verification*

Pass 2 · Downgrade
 $g_\downarrow(q, r, \hat{y}) \in \{0, 1\}$
applies if $\hat{y}_1 \in \{\text{CR}, \text{CNR}\}$
one vs. multiple interpretations?
 $g_\downarrow = 1 \Rightarrow y^\downarrow = \text{AMB}$

CR / CNR

29 pairs

$g_\downarrow = 1$

3 pairs

AMB
class

check ↑

120 pairs

Pass 3 · Upgrade
 $g_\uparrow(q, r) \in \{0, 1\}$
applies if label = AMB
first substantive sentence commits?
 $g_\uparrow = 1 \Rightarrow \hat{y} = \text{CR}$

$\hat{y} = \text{CR}$ if $g_\uparrow = 1$,
else AMB

both corrections pass through AMB (edge counts: routing of the 308 dev pairs, GLM-4.7 full config)

$g_\downarrow = 0$: stays CR/CNR (26)